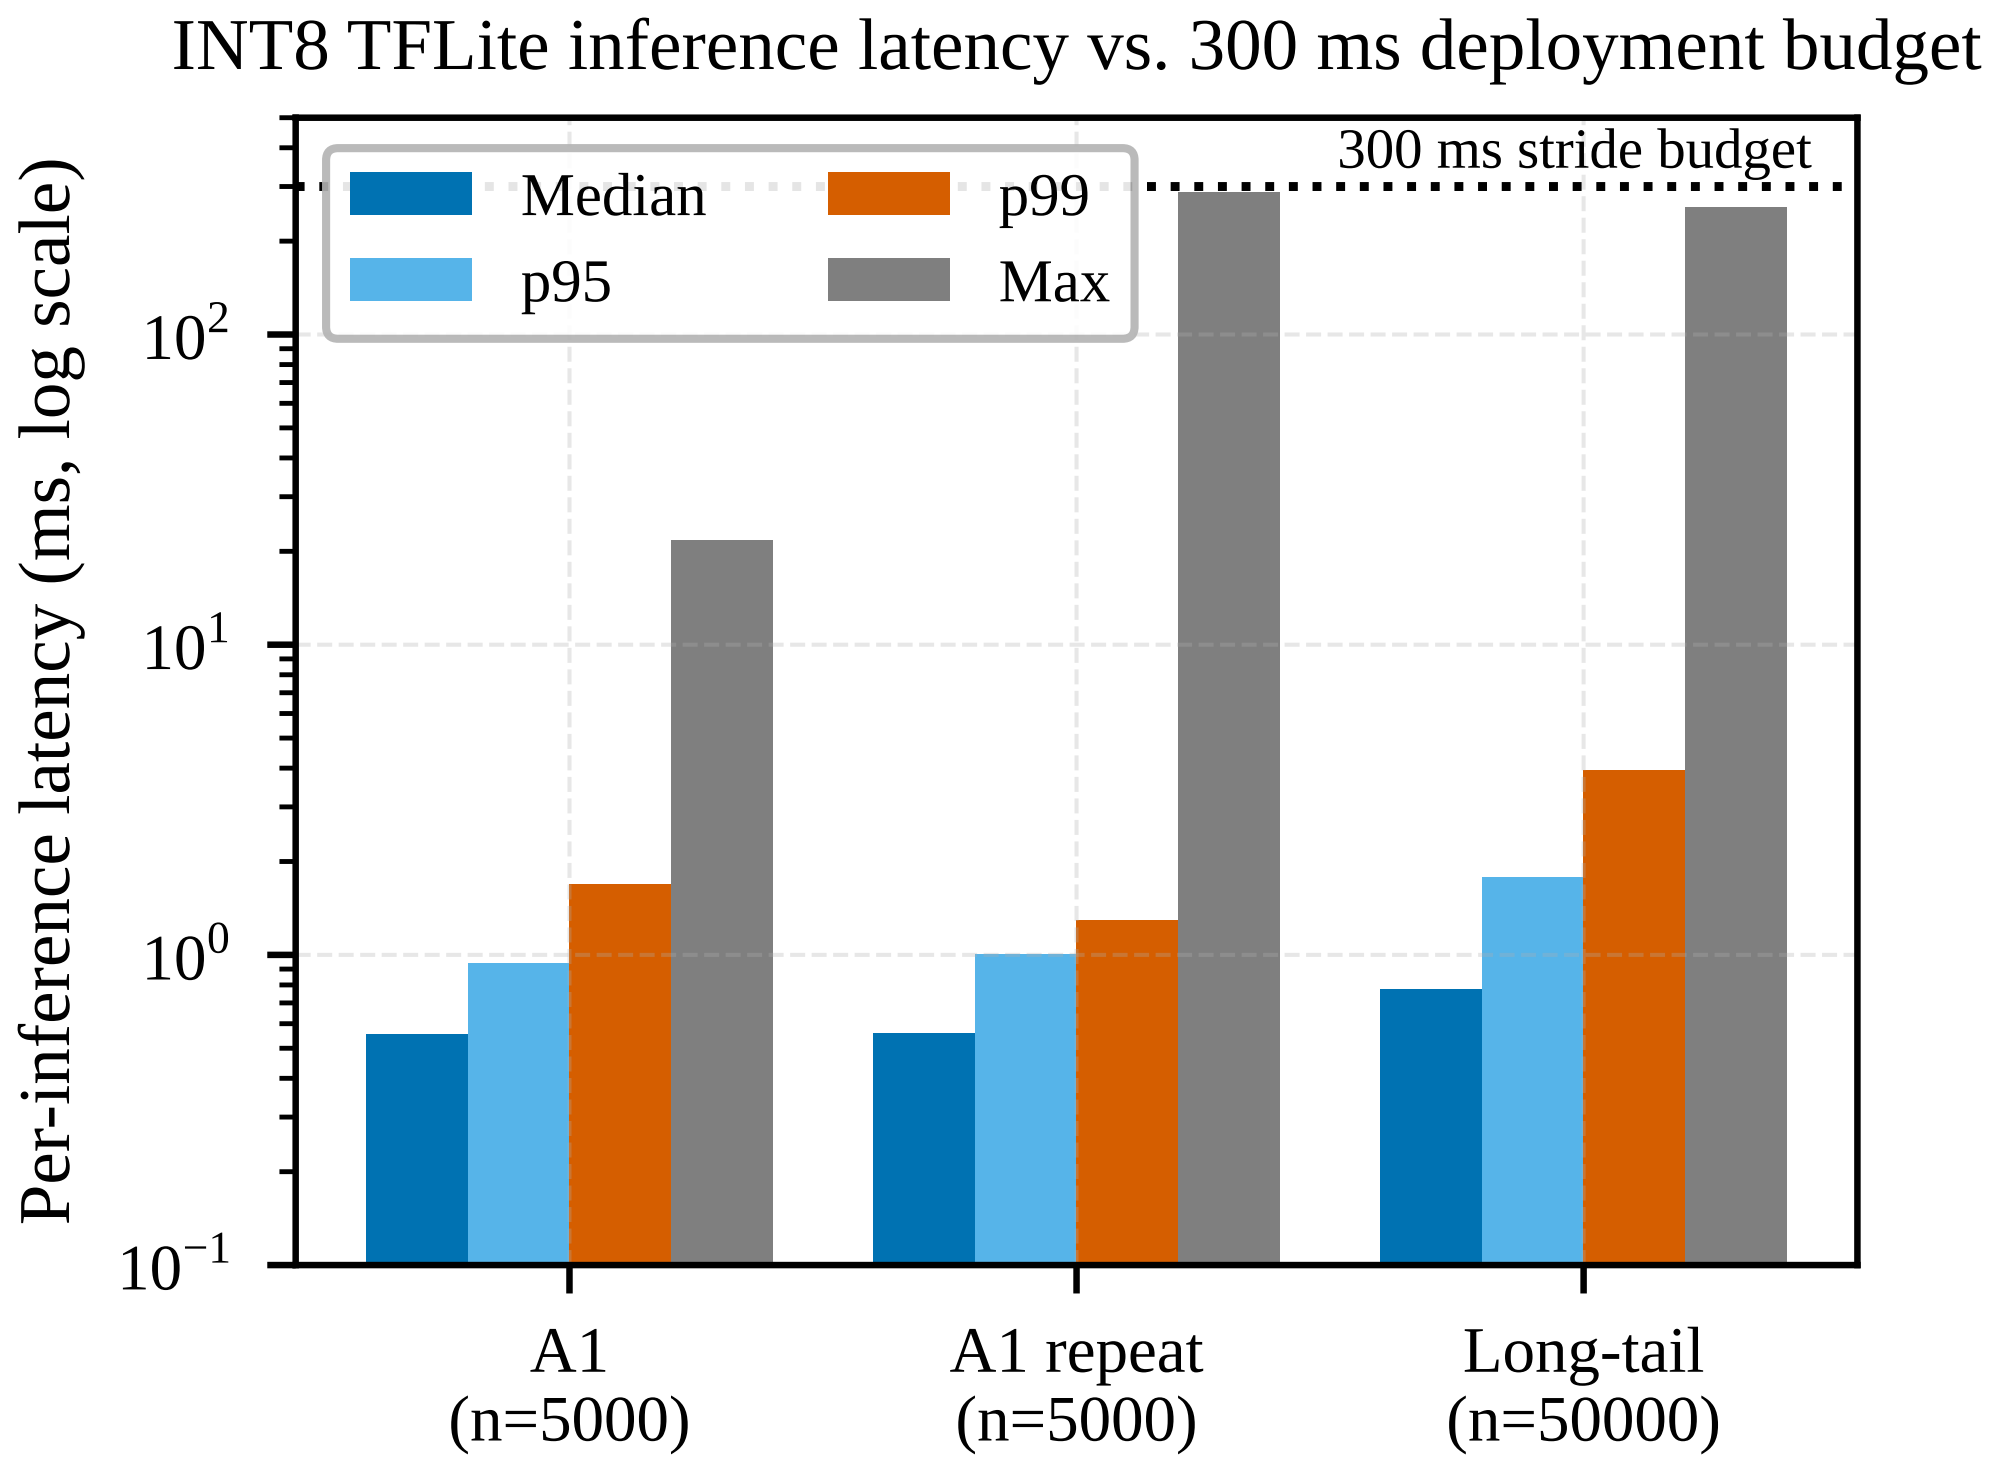
INT8 TFLite inference latency vs. 300 ms deployment budget
Median
p95
p99
Max
300 ms stride budget
Per-inference latency (ms, log scale)
10^{2}
10^{1}
10^{0}
10^{-1}
A1
(n=5000)
A1 repeat
(n=5000)
Long-tail
(n=50000)