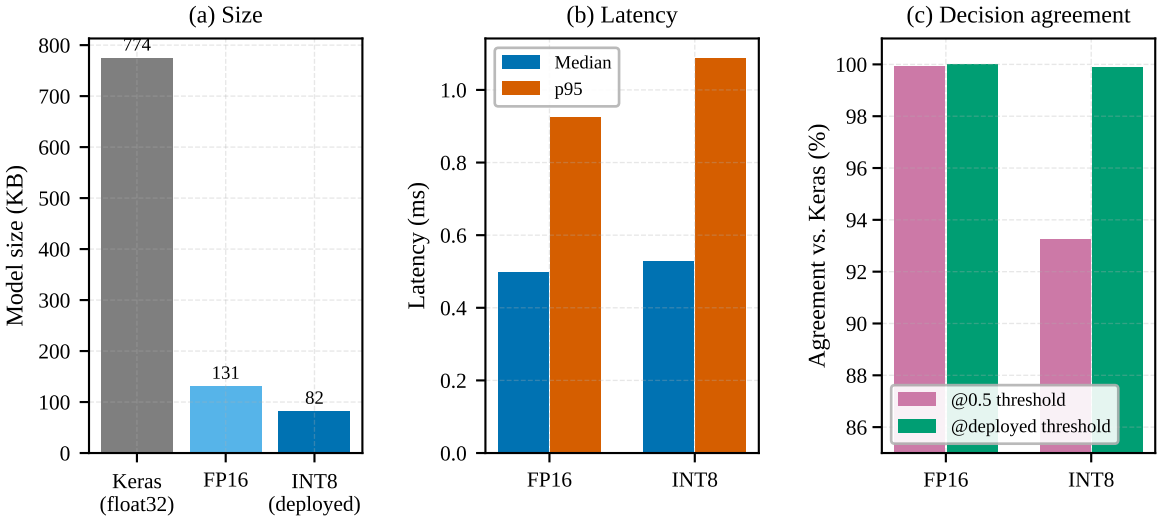
(a) Size
Model size (KB)
800
700
600
500
400
300
200
100
0
774
131
82
Keras (float32)
FP16
INT8 (deployed)
(b) Latency
Latency (ms)
1.0
0.8
0.6
0.4
0.2
0.0
Median
p95
FP16
INT8
(c) Decision agreement
Agreement vs. Keras (%)
100
98
96
94
92
90
88
86
@0.5 threshold
@deployed threshold
FP16
INT8